

# **Photogrammetry & Robotics Lab**

## **Machine Learning for Robotics and Computer Vision**

### **Introduction**

**Jens Behley**

---

# What is Machine Learning?

- Classical definition by Tom Mitchell:

A computer program is said to **learn** from *experience*  $E$  with respect to some class of *task*  $T$ , and *performance measure*  $P$ , if its performance at tasks in  $T$ , as measured by  $P$ , improves with experience  $E$ .

From: Tom M. Mitchell, Machine Learning, McGraw-Hill, 1997

# But what is Machine Learning?

- Rather broad definition that captures the many faces of Machine Learning (ML)
- **Goal:** We want to design algorithms that **automatically** extract valuable information from **data** via adapting a **model**
- Adaptation of the model based on data is what is called **learning**

# Examples from Everyday Life

- Many things are powered by ML:
  - Spam filter for your emails
  - Auto correct on smart phones
  - Recommendations on streaming services
  - Speech recognition in personal assistants
  - ...
- Often tasks where it's hard to write a program by fixed rules
- ML solves this by learning from a large set of examples (= experiences)

# Examples in Robotics & Computer Vision



- Perception in self-driving cars
- Semantic interpretation of images

# Contents of the Course

- **Part I:** Traditional ML methods
  - Basics, ML terminology, general ML models for classification, regression & clustering
- **Part II:** Deep Learning for Vision Tasks
  - Convolutional Neural Networks (CNN), learning CNNs, current research topics

# People



Jens Behley



Lucas Nunes

# Lecture, Exercises, and Exam

- Lectures as **video** recordings
- Tutorials & questions via **Zoom**
- **eCampus** website for further information
- Homework assignments
- **Deadlines:** see eCampus deadlines
- **Oral** exam (most likely via Zoom)




# Homework Assignments

- **Exam Admission:** at least **50%** of the points overall
- Coding (Python) is an essential part of the homework assignments
- Assignment/submission via **eCampus**
- No plagiarism! Copying of solutions is not accepted (**zero tolerance policy**)

# Books

- Most of the contents are based on two freely available books



<https://mml-book.com>

[MML]

[MML Chap. 9.3]



<http://www.computer-visionmodels.com/>

[CVMLI]

# Doing a better job...



Source: <https://xkcd.com/1838/>


# High-level overview

- **Main components** of ML algorithm
  1. **Data**
  2. **Model**
  3. **Learning**
- Here, high-level overview of the concepts
- In next lectures, we will dive into specific models for specific tasks and cover learning

# Types of Learning

- **Supervised** Learning
  - Examples with target values (= **labels**) available
  - Targets can be **continuous** or **discrete**
  - **Goal:** Learn predictors for inferring the target value for unseen data
- **Unsupervised** Learning
  - Examples without explicit targets (unlabeled)
  - **Goal:** Analysis of the data or learn generative models or representations of the data

# Regression vs. Classification



$\mathbf{x}$	$y$
(0.8, 0.2, 0.3, 0.9)	1
(0.3, 0.2, 0.3, 0.4)	2
(0.1, 0.9, 0.4, 0.1)	1
(0.5, 0.1, 0.2, 0.2)	3
(0.1, 0.2, 0.8, 0.7)	?

## ▪ Regression

- Input: example  $\mathbf{x} \in \mathbb{R}^D$
- Output: target values  $y \in \mathbb{R}$  (continuous)

## ▪ Classification

- Input: example  $\mathbf{x} \in \mathbb{R}^D$
- Output: class  $y \in \{1, \dots, k\} = \mathcal{Y}$  (discrete)

# Data

# Where do we get the data?

- **Research:** (Benchmark) datasets available that provide data in the right format
  - Advantage: authors of the dataset ensured that data is ready to go
  - Disadvantage: Specific domain & dataset biases, may not fit your task
- Getting data for new tasks & domains
  - Retrieve from databases or crawl internet
  - Collect/record new data (robot, camera, ...)



# Unstructured Data



- Data from multiple sources is usually unstructured
- Data can have different types of fields
  - Categorical values, e.g., "red", "green", ...
  - Strings, e.g., names, ...
  - Missing values, e.g., " ", "NULL", ...
  - Misspelled values
  - etc.

# Data Wrangling

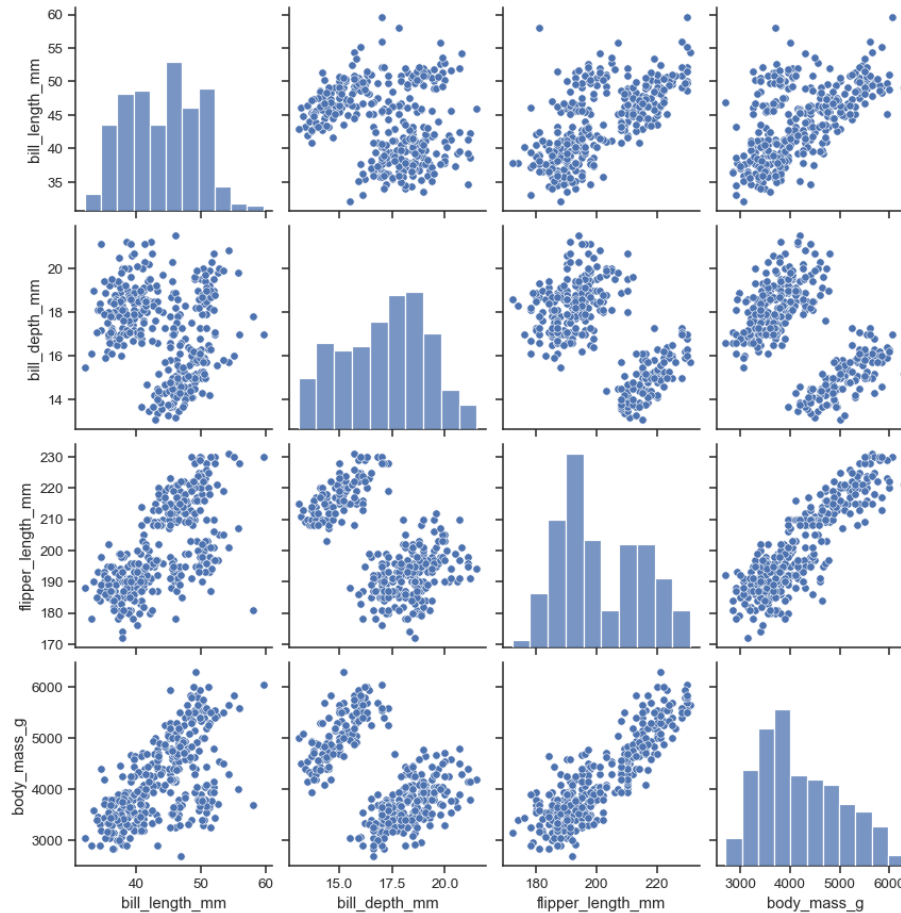
- Process of transforming raw data into a usable format
- Main task: cleaning and converting values
- Handling missing values
  - Replace with 0 (= "ignore feature")
  - Impute values, e.g., mean of other values, regression, ...
  - Drop example (worst option)

# Visualize your data!

- Summary statistics are not enough!
- Let's say you have a dataset  $\mathcal{X}, \mathbf{x}_n \in \mathbb{R}^2$ , with following statistics (mean and standard deviation):  
 $\mu = (54.28, 47.83), \sigma = (16.76, 26.93)$



# Plotting multi-dimensional Data



Source: seaborn.pairplot

- Plot pairs that might show dependencies and relations

# Data in Supervised Learning

- Examples in supervised learning given as set of tuples

$$\mathcal{X} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n), \dots, (\mathbf{x}_N, y_N)\}$$

- Data vectors  $\mathbf{x}_n \in \mathbb{R}^D$  called **feature vector**
- Entry  $x_d \in \mathbb{R}$  of data vector
$$\mathbf{x}_n = (x_1, \dots, x_d, \dots, x_D)$$
is called **feature**, **attribute**, or covariate
- $y$  is called **label**, target, response, or annotation
- Data used for learning is called **training set**  $\mathcal{X}_{\text{train}}$

# Example: Apples vs. Pears

- Let's say we want to classify fruit images
- **Features:** "roundness" and "redness"
- Image from which we determine the following features
  - $x_1 \in [0, 1]$  = ratio of radii of fitted ellipse radii
  - $x_2 \in [0, 1]$  = avg. of red values of fruit pixels
- **Output:**
  - $y = \{0, 1\} \rightarrow \{\text{Apple, Pear}\}$
  - **→ binary classification** problem

# Example: Apple vs. Pears

- Gathered images, curated them, extracted features




© Abhijit Tembhekar

→ (0.95, 0.96)



© I, NibbiP

→ (0.15, 0.5)



© PiccoloNamek

→ (0.96, 0.1)

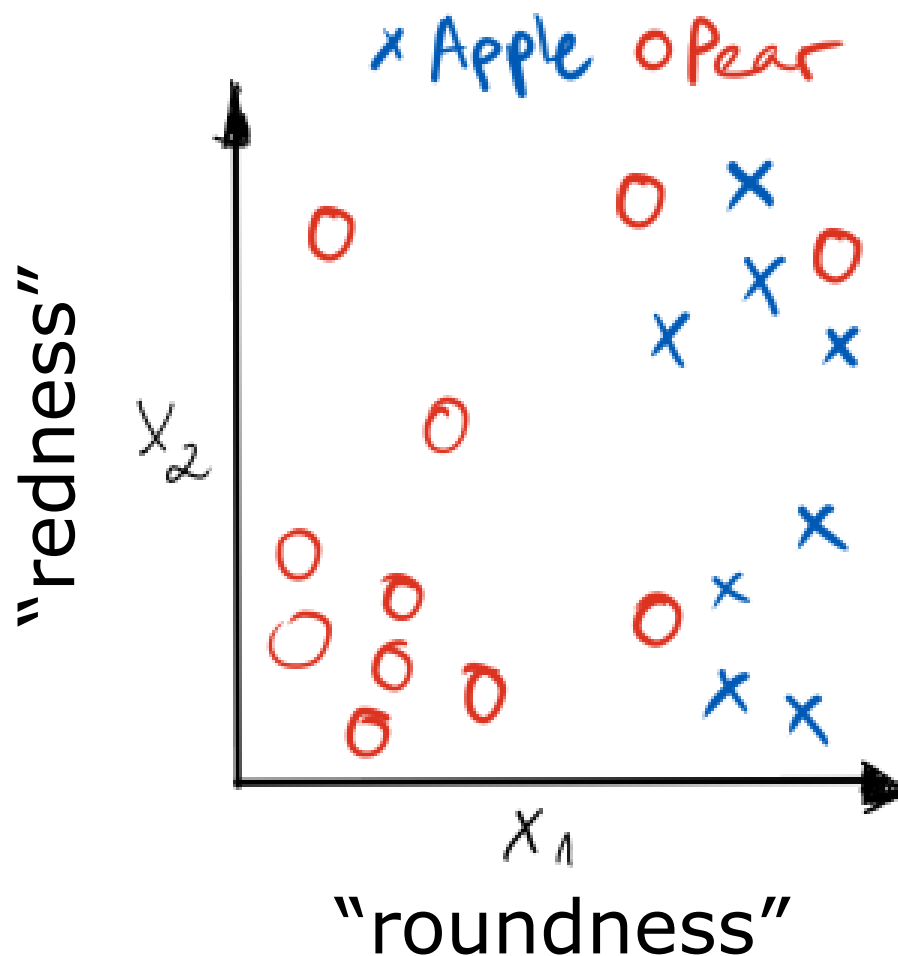


© CDC

→ (0.95, 0.9)

# Example: Apple vs. Pears

- We get the following **feature space** where each point corresponds to an example:





# Model

# Function vs. Probabilistic Model

- Different views on the role of the model as a predictor
- Models as **functions**  $f : \mathbb{R}^D \mapsto \mathbb{R}$  producing an output  $y$  for given input  $\mathbf{x} \in \mathbb{R}^D$ :
$$y = f(x)$$
- Models as **probability distribution**  $P(y|\mathbf{x})$ :
$$\hat{y} = \arg \max_{c \in \mathcal{Y}} P(y = c|\mathbf{x})$$
- Advantage: Get confidence in prediction

# Parameters & Hyperparameters

- The model is specified by adjustable **parameters**  $\theta$  that are learned from data
- Amount of model parameters determine the **capacity** of the model
- **Hyperparameters** = Parameters of the model that are not learned from data
- Hyperparameters are selected in advance

# Example: Apple vs. Pears (cont)

## k-Nearest Neighbor Classifier


- Example: **k-Nearest Neighbor Classifier**
- Let  $\mathcal{N}_k(\mathbf{x})$  the k-nearest neighbors of  $\mathbf{x}$  from  $(\mathbf{x}_n, y) \in \mathcal{X}_{\text{train}}$
- Then  $P(y = c|\mathbf{x})$  is given by

$$P(y = c|\mathbf{x}) = \frac{1}{k} |\{(x_k, y_k) \in \mathcal{N}_k(\mathbf{x}) \mid y_k = c\}|$$

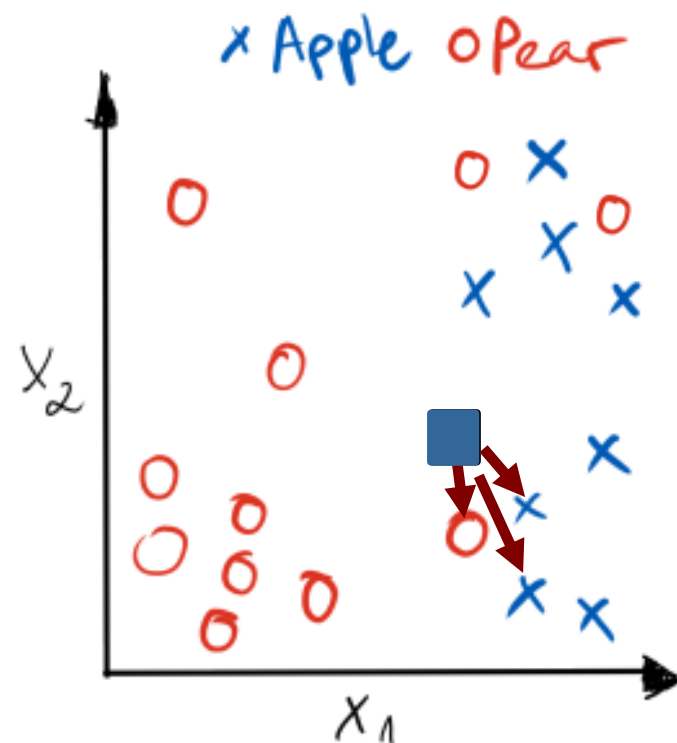
- **Hyperparameter:** number of neighbors k

# Example: Apple vs. Pears (cont)

## k-Nearest Neighbor Classifier



$k=1$



$k=3$

- Depending on  $k$  different outputs

# Learning

# Learning the model

- Learning is adapting the **parameters**  $\theta$  of the model given the training set  $\mathcal{X}_{\text{train}}$
- In most cases its maximizing the likelihood

$$P(y_1, \dots, y_N | \mathbf{x}_1, \dots, \mathbf{x}_n, \theta)$$

- Assumption: set of examples are **independent and identically distributed**
- Thus, learning is an optimization problem:

$$\theta^* = \arg \max_{\theta} \prod_{n=1}^N P(y_n | \mathbf{x}_n)$$

# Example: Apple vs. Pears (cont)

## k-Nearest Neighbor Classifier

- Learning the k-Nearest Neighbors classifier is simple → store the training set  $\mathcal{X}_{\text{train}}$
- How well does the model perform?
- How to choose hyperparameter k?



# Evaluation

# Measuring Performance

- **Goal:** Learn models that performs well on unseen data → generalization
- How do we measure performance? Metrics!
- **Training error** = how well do we determine target values  $y_n$  for each  $x_n$
- Other metrics of interest possible
  
- **Problem:** training error just tells us performance on seen data; should be ideally close to zero

# Training and Test set

- **Solution:** Evaluate model performance on part of the data not used for training!
- Data used for learning, called **training**, is called **training set**  $\mathcal{X}_{\text{train}}$
- Data used for evaluating trained model is **test set**  $\mathcal{X}_{\text{test}}$
- It holds  $\mathcal{X}_{\text{train}} \cap \mathcal{X}_{\text{test}} = \emptyset$  !
- Test set **not** used for determining the parameters!
- Test set should be only used seldomly!

# Human in the Loop Phenomenon


- Test set is our *estimate* how algorithm works on **unseen** data
- But you might be tempted to just “look” at the test data and select hyperparameters or parameters that “look” good on the test data
- Overfitting to test data and biased result possible → no insight about generalization
- “Good looking” model might fail miserably in real world application!

# Validation set

- **Selecting hyperparameters** should happen on a so-called **validation set**
- Validation set is usually part from the training set
- But: validation set is also not used for training
  - With fixed hyperparameters sometimes still used for re-training the model)
- **Validation set  $\neq$  test set!**

# Example: Apple and Pears

- Selecting  $k$  such that validation error is minimized.
- Error =  $\frac{|\{\mathbf{x}_n \in \mathcal{X} | \hat{y} \neq y_n\}|}{|\mathcal{X}|}$  with  $\hat{y} = \arg \max_{c \in \mathcal{Y}} P(y = c | \mathbf{x})$



# Again, don't fool yourself!

- Tuning on test set will lead to wrong conclusions!
- Good performance on your test set will not translate to good performance in real world if you select parameters on the test set!
- All hyperparameter choices should only happen on the validation set

# Summary

- Covered high-level overview of ML algorithms
- Introduced some ML terminology
- High-level overview of the ML pipeline
  
- Next lectures:
  - More on regression and classification models
  - Mostly, details on more complex models and learning them



**See you next week!**